

1003 Hw 7: Computation Graphs, Backpropagation, and Neural Networks

rjl407

May 2019

1 Introduction

2 Computation Graph Framework

3 Ridge Regression

1. Complete the class L2NormPenaltyNode in nodes.py.

```
1 class L2NormPenaltyNode(object):
2     """ Node computing l2_reg * ||w||^2 for scalars l2_reg and
3         vector w"""
4     def __init__(self, l2_reg, w, node_name):
5         """
6         Parameters:
7         l2_reg: a scalar value >=0 (not a node)
8         w: a node for which w.out is a numpy vector
9         node_name: node's name (a string)
10        """
11        self.node_name = node_name
12        self.out = None
13        self.d_out = None
14        self.l2_reg = l2_reg
15        self.w = w
16
17        # TODO
18
19    def forward(self):
20        self.out = self.l2_reg * np.dot(self.w.out, self.w.out)
21
22        self.d_out = np.zeros(self.out.shape)
23        return self.out
24
25    def backward(self):
26        d_w = self.d_out * 2 * self.l2_reg * self.w.out
27        self.w.d_out += d_w
28        return self.d_out
29
30    def get_predecessors(self):
31        return [self.w]
```

2. Complete the class SumNode in nodes.py.

```

1 class SumNode(object):
2     """ Node computing a + b, for numpy arrays a and b """
3     def __init__(self, a, b, node_name):
4         """
5         Parameters:
6         a: node for which a.out is a numpy array
7         b: node for which b.out is a numpy array of the same
8         shape as a
9         node_name: node's name (a string)
10        """
11        # TODO
12        self.node_name = node_name
13        self.out = None
14        self.d_out = None
15        self.a = a
16        self.b = b
17
18    def forward(self):
19        self.out = self.a.out + self.b.out
20        self.d_out = np.zeros(self.out.shape)
21        return self.out
22
23    def backward(self):
24        d_a = self.d_out
25        d_b = self.d_out
26        self.a.d_out += d_a
27        self.b.d_out += d_b
28        return self.d_out
29
30    def get_predecessors(self):
31        return [self.a, self.b]

```

3. Implement ridge regression with w regularized and b unregularized. Do this by completing the `__init__` method in the `ridge_regression.py`, using the classes created above. When complete, you should be able to pass the tests in `ridge_regression.t.py`. Report the average square error on the **training** set for the parameter settings given in the `main()` function.

Answer:

`l2reg = 1`, Ave training loss: 0.19972371048423804

`l2reg = 0`, Ave training loss: 0.030071862072324693

```

1 class RidgeRegression(BaseEstimator, RegressorMixin):
2     """ Ridge regression with computation graph """
3     def __init__(self, l2_reg=1, step_size=.005,
4         max_num_epochs = 5000):
5         self.max_num_epochs = max_num_epochs
6         self.step_size = step_size
7
8         # Build computation graph
9         self.x = nodes.ValueNode(node_name="x") # to hold a
10        vector input
11        self.y = nodes.ValueNode(node_name="y") # to hold a
12        scalar response

```

```


10     self.w = nodes.ValueNode(node_name="w") # to hold the
        parameter vector
11     self.b = nodes.ValueNode(node_name="b") # to hold the
        bias parameter (scalar)
12     self.prediction = nodes.VectorScalarAffineNode(x=self.
        x, w=self.w, b=self.b,
13                                                     node_name="
        prediction")
14     # TODO
15     self.square_loss = nodes.SquaredL2DistanceNode(a=self.
        prediction, b=self.y,
16                                                     node_name="
        square loss")
17     self.l2_norm = nodes.L2NormPenaltyNode(l2_reg, self.w,
        node_name="l2_norm")
18     self.objective = nodes.SumNode(self.square_loss, self.
        l2_norm, node_name="objective")
19     self.inputs = [self.x]
20     self.outcomes = [self.y]
21     self.parameters = [self.w, self.b]
22     self.graph = graph.ComputationGraphFunction(self.
        inputs, self.outcomes,
23
        self.parameters, self.prediction,
24
        self.objective)
25
26     def fit(self, X, y):
27         num_instances, num_ftrs = X.shape
28         y = y.reshape(-1)
29
30         init_parameter_values = {"w": np.zeros(num_ftrs), "b":
        np.array(0.0)}
31         self.graph.set_parameters(init_parameter_values)
32
33         for epoch in range(self.max_num_epochs):
34             shuffle = np.random.permutation(num_instances)
35             epoch_obj_tot = 0.0
36             for j in shuffle:
37                 obj, grads = self.graph.get_gradients(
        input_values = {"x": X[j]},
38
        outcome_values = {"y": y[j]})
39                 epoch_obj_tot += obj
40                 # Take step in negative gradient direction
41                 steps = {}
42                 for param_name in grads:
43                     steps[param_name] = -self.step_size *
        grads[param_name]
44                 self.graph.increment_parameters(steps)
45
46                 if epoch % 50 == 0:
47                     train_loss = sum((y - self.predict(X,y)) **2)/
        num_instances
48                     print("Epoch ", epoch, ": Ave objective=",
        epoch_obj_tot/num_instances, " Ave training loss: ",
        train_loss)

```

```

49
50 def predict(self, X, y=None):
51     try:
52         getattr(self, "graph")
53     except AttributeError:
54         raise RuntimeError("You must train classifier
before predicting data!")
55
56     num_instances = X.shape[0]
57     preds = np.zeros(num_instances)
58     for j in range(num_instances):
59         preds[j] = self.graph.get_prediction(input_values
= {"x": X[j]})
60
61     return preds

```



4. [Optional] Create a new implementation of ridge regression that supports efficient minibatching. You will replace the the ValueNode x , which contains a vector, with a ValueNode X , which contains a matrix. The convention is that the first dimension indexes examples and the second is features (as we have always done). Many of the nodes will have to be adapted to this use case. Demonstrate its use and speedup.

$$4.1.1.1 \quad \frac{\partial y_r}{\partial w_{ij}} = \begin{cases} x_j & , r=i \\ 0 & , r \neq i \end{cases}$$

$$\frac{\partial J}{\partial w_{ij}} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial w_{ij}} = \frac{\partial J}{\partial y_i} \cdot x_j$$

$$4.1.1.2 \quad \frac{\partial J}{\partial W} = \frac{\partial J}{\partial y} \otimes x = \frac{\partial J}{\partial y} x^T$$

$$4.1.1.3 \quad \frac{\partial J}{\partial x} = \frac{\partial J}{\partial y} \cdot \frac{\partial y}{\partial x} = W^T \frac{\partial J}{\partial x}$$

$$4.1.1.4 \quad \frac{\partial J}{\partial b} = \frac{\partial J}{\partial y} \cdot \frac{\partial y}{\partial b} = \frac{\partial J}{\partial y}$$

$$4.1.2. \quad \frac{\partial J}{\partial A} = \frac{\partial J}{\partial S} \cdot \frac{\partial S}{\partial A} = \frac{\partial J}{\partial S} \odot \sigma'(A)$$

4 Multilayer Perceptron

4.1

4.2 MLP Implementation

1. Complete the class `AffineNode` in `nodes.py`. Be sure to propagate the gradient with respect to x as well, since when we stack these layers, x will itself be the output of another node that depends on our optimization parameters.

```
1 class AffineNode(object):
2     """Node implementing affine transformation (W,x,b)→Wx+b,
3     where W is a matrix,
4     and x and b are vectors
5     Parameters:
6     W: node for which W.out is a numpy array of shape (m,d
7     )
8     x: node for which x.out is a numpy array of shape (d)
9     b: node for which b.out is a numpy array of shape (m)
10    (i.e. vector of length m)
11    """
12
13    ## TODO
14    def __init__(self, W, x, b, node_name):
15        self.node_name = node_name
16        self.out = None
17        self.d_out = None
18        self.x = x
19        self.W = W
20        self.b = b
21
22    def forward(self):
23        self.out = np.dot(self.W.out, self.x.out) + self.b.out
24        self.d_out = np.zeros(self.out.shape)
25        return self.out
26
27    def backward(self):
28        d_b = self.d_out
29        d_x = np.matmul(np.transpose(self.W.out), self.d_out)
30        d_W = np.matmul(self.d_out.reshape(-1, 1), np.
31        transpose(self.x.out.reshape(-1,1)))
32        self.b.d_out += d_b
33        self.x.d_out += d_x
34        self.W.d_out += d_W
35        return self.d_out
36
37    def get_predecessors(self):
38        return [self.x, self.W, self.b]
```

2. Complete the class `TanhNode` in `nodes.py`. As you'll recall, $\frac{d}{dx} \tanh(x) = 1 - \tanh^2 x$. Note that in the forward pass, we'll already have computed \tanh of the input and stored it in `self.out`. So make sure to use `self.out` and not recalculate it in the backward pass.

```

1 class TanhNode(object):
2     """Node tanh(a), where tanh is applied elementwise to the
3     array a
4     Parameters:
5     a: node for which a.out is a numpy array
6     """
7     ## TODO
8     def __init__(self, a, node_name):
9         self.node_name = node_name
10        self.out = None
11        self.d_out = None
12        self.a = a
13
14    def forward(self):
15        self.out = np.tanh(self.a.out)
16        self.d_out = np.zeros(self.out.shape)
17        return self.out
18
19    def backward(self):
20        d_a = (1 - self.out * self.out) * self.d_out
21        self.a.d_out += d_a
22        return self.d_out
23
24    def get_predecessors(self):
25        return [self.a]

```

3. Implement an MLP by completing the skeleton code in `mlp_regression.py`, and making use of the nodes above. Your code should pass the tests provided in `mlp_regression.t.py`. Note that to break the symmetry of the problem, we initialize our weights to small random values, rather than all zeros, as we often do for convex optimization problems. Run the MLP for the two settings given in the `main()` function and report the average **training** error. Note that with an MLP, we can take the original scalar as input, in the hopes that it will learn nonlinear features on its own, using the hidden layers. In practice, it is quite challenging to get such a neural network to fit as well as one where we provide features.

Traning errors:

no features: 0.2190364498138465

with features: 0.027023495833501273

```

1 class MLPRegression(BaseEstimator, RegressorMixin):
2     """ MLP regression with computation graph """
3     def __init__(self, num_hidden_units=10, step_size=.005,
4     init_param_scale=0.01, max_num_epochs = 5000):
5         self.num_hidden_units = num_hidden_units
6         self.init_param_scale = 0.01
7         self.max_num_epochs = max_num_epochs
8         self.step_size = step_size
9
10        # Build computation graph
11        self.x = nodes.ValueNode(node_name="x") # to hold a
12        vector input

```

```

11     self.y = nodes.ValueNode(node_name="y") # to hold a
scalar response
12     ##TODO
13     self.W1 = nodes.ValueNode(node_name="W1")
14     self.b1 = nodes.ValueNode(node_name="b1")
15     self.W2 = nodes.ValueNode(node_name="w2")
16     self.b2 = nodes.ValueNode(node_name="b2")
17     self.L = nodes.AffineNode(self.W1, self.x, self.b1,
node_name="L")
18     self.h = nodes.TanhNode(self.L, node_name="h")
19     self.prediction = nodes.VectorScalarAffineNode(self.h,
self.W2, self.b2, node_name="prediction")
20     self.objective = nodes.SquaredL2DistanceNode(a=self.
prediction, b=self.y,
21                                                     node_name="
objective")
22     self.inputs = [self.x]
23     self.outcomes = [self.y]
24     self.parameters = [self.W1, self.b1, self.W2, self.b2]
25     self.graph = graph.ComputationGraphFunction(self.
inputs, self.outcomes,
26
self.parameters, self.prediction,
27
self.objective)
28
29     def fit(self, X, y):
30         num_instances, num_ftrs = X.shape
31         y = y.reshape(-1)
32
33         ## TODO: Initialize parameters (small random numbers
— not all 0, to break symmetry )
34         s = self.init_param_scale
35         init_values = {
36             "W1": s * np.random.normal(0,1,size=(self.
num_hidden_units, num_ftrs)), #
gaussian01
37             "b1": np.zeros(self.num_hidden_units),
38             "w2": s * np.random.normal(0,1,self.
num_hidden_units), # gaussian
39             "b2": np.array(0.0)
40         }
41
42         self.graph.set_parameters(init_values)
43
44         for epoch in range(self.max_num_epochs):
45             shuffle = np.random.permutation(num_instances)
46             epoch_obj_tot = 0.0
47             for j in shuffle:
48                 obj, grads = self.graph.get_gradients(
input_values = {"x": X[j]},
49
outcome_values = {"y": y[j]})
50                 #print(obj)
51                 epoch_obj_tot += obj
52                 # Take step in negative gradient direction
53                 steps = {}

```



```

54         for param_name in grads:
55             steps[param_name] = -self.step_size *
grads[param_name]
56             self.graph.increment_parameters(steps)
57
58         if epoch % 50 == 0:
59             train_loss = sum((y - self.predict(X,y)) **2)/
num_instances
60             print("Epoch ",epoch," : Ave objective=",
epoch_obj_tot/num_instances," Ave training loss: ",
train_loss)
61
62     def predict(self, X, y=None):
63         try:
64             getattr(self, "graph")
65         except AttributeError:
66             raise RuntimeError("You must train classifier
before predicting data!")
67
68         num_instances = X.shape[0]
69         preds = np.zeros(num_instances)
70         for j in range(num_instances):
71             preds[j] = self.graph.get_prediction(input_values
= {"x":X[j]})
72
73         return preds

```

